

CSCE 3600: Systems Programming

Major Assignment 1 – The Shell and System Calls

Due: 11:59 PM on Friday, April 6, 2018

COLLABORATION

You should complete this assignment as a group assignment with the other members of your group. *Each group should have 3 or 4 members, no more, no less.* Submit only ONE program per group. Also, make sure that you list the names of all group members who participated in this assignment in order for each to get credit.

BACKGROUND

A shell provides a command-line interface for users. It interprets user commands and executes them. Some shells provide simple scripting terms, such as `if` or `while`, and allow users to make a program that facilitates their computing environment. Under the hood, a shell is just another user program. The file `/bin/bash` is an executable file for the bash shell. The only thing special about your login shell is that it is listed in your login record so that `/bin/login` (i.e., the program that prompts you for your password) knows what program to start when you log in. If you run `"cat /etc/passwd"`, you will see the login records of the machine and the login shell program listed as the last field.

PROGRAM DESCRIPTION

In this assignment, you will implement:

1. a command line interpreter, or shell;

Your shell should read the line from standard input (i.e., interactive mode) or a file (i.e., batch mode), parse the line with command and arguments, execute the command with arguments, and then prompt for more input (i.e., the shell prompt) when it has finished.

a. Interactive Mode

In interactive mode, you will display a prompt (any string of your choosing) and the user of the shell will type in a command at the prompt.

b. Batch Mode

In batch mode, your shell is started by specifying a batch file on its command line. The batch file contains the list of commands that should be executed. In batch mode, you should not display a prompt, but you should echo each line you read from the batch file back to the user before executing it.

You will need to use the `fork()` and `exec()` family of system calls. You may not use the `system()` system call as it simply invokes the system's `/bin/bash` shell to do all of the work.

You may assume that arguments are separated by whitespace. You do not have to deal with special characters such as `'`, `"`, `\`, etc. However, you will need to handle the redirection operators (`<` and `>`) and the pipeline operator (`|`).

Each line (either the batch file or typed at the prompt) may contain multiple commands separated with the semicolon (`;`) character. Each command separated by a `;` should be run simultaneously, or concurrently. Note that this is different behavior than standard Linux shells that run these commands one at a time, in order. The shell should not print the next prompt or take more input until all of these commands have finished executing (the `wait()` or `waitpid()` system calls may be useful here).

You may assume that the command-line a user types is not longer than 512 bytes (including the `'\n'`), but you should not assume that there is any restriction on the number of arguments to a given command.

2. the following built-in commands;

Every shell needs to support a number of built-in commands, which are functions in the shell itself, not external programs. Shells directly make system calls to execute built-in commands, instead of forking a child process to handle them.

- Add a new built-in `exit` command that exits from the shell itself with the `exit()` system call. It is not to be executed like other programs the user types in. If the `exit` command is on the same line with other commands, you should ensure that the other commands execute (and finish) before you exit your shell.

These are all valid examples for quitting the shell:

```
prompt> exit
prompt> exit; cat file1
prompt> cat file1; exit
```

- Add a new built-in `pwd` command that prints the current working directory to standard output. You may need to invoke the `getcwd()` system call.
- Add a new built-in `cd` command that accepts one optional argument, a directory path, and changes the current working directory to that directory. If no argument is passed, the command will change the current working directory to the user's `HOME` directory. You may need to invoke the `chdir()` system call.

3. extend your shell with I/O redirection;

When you start a command, there are always three default file streams open: `stdin` (maps input from the keyboard by default), `stdout` (maps output to the terminal by default), and `stderr` (maps error messages to the terminal by default). These and other open files may be redirected, or mapped, to files or devices that users specify.

Modify your shell so that it supports redirecting `stdin` and `stdout` to files. You

do not need to support redirection for shell built-in commands (i.e., `exit`, `cd`, and `pwd`). You do not need to support `stderr` redirection or appending to files (e.g., `cmd3 >> out.txt`). You may assume that there will always be spaces around the special characters `<` and `>`. Be aware that the "`< file`" or "`> file`" are not passed as arguments to the program.

Some redirection examples include:

```
$ cmd1 < in.txt
```

executes `cmd1`, using `in.txt` as the source of input, instead of the keyboard.

```
$ cmd2 > out.txt
```

executes `cmd2` and places the output to file `out.txt`.

You will need to understand Linux file descriptors and use the `open()`, `close()`, and `dup()` family of system calls. Your shell should be able to handle cases of double redirection (i.e., both input and output).

4. extend your shell with pipelining; and

The command

```
$ cmd1 | cmd2 | cmd3
```

connects the standard output of `cmd1` to the standard input of `cmd2`, and again connects the standard output of `cmd2` to the standard input of `cmd3` using the pipeline operator `|`. An example usage is

```
$ sort < file.txt | uniq | wc
```

counts the number of unique lines in `file.txt`. Without pipes, you would have to use three commands and two intermediate files to count the unique lines in a file.

You will need to use the `pipe()` system call. Your shell should be able to handle up to three commands chained together with the pipeline operator (i.e., your shell should support up to two pipes pipelined together). As seen in the usage example, your shell should also handle combinations of redirection and pipelining when they can be combined.

Your shell does not need to handle built-in commands implemented above (i.e., `cd`, `exit`, and `pwd`) in pipeline.

5. support of signal handling and terminal control.

Many shells allow you to stop or pause processes with special keystrokes, such as `Ctrl-C` or `Ctrl-Z`, that work by sending signals to the shell's subprocesses. If you try these keystrokes in your shell, the signals would be sent directly to the shell process itself. This is not what we want since, for example, attempting to `Ctrl-Z` a subprocess of your shell will also stop the shell itself. Instead, we want to have the signals affect only the subprocesses that our shell creates. To help

you accomplish this, you might find the following helpful:

a. Process Groups

Every process has a unique process ID (i.e., `pid`). Every process also has a possibly non-unique process group ID (i.e., `pgid`) which, by default, is the same as the `pgid` of its parent process. Processes can get and set their process group ID with the system calls `getpgid()`, `setpgid()`, `getpgrp()`, or `setpgrp()`.

Keep in mind that, when your shell starts a new program, that program might require multiple processes to function correctly. All of these processes will inherit the same process group ID of the original process. So, it may be a good idea to put each shell process into its own process group for simplicity. When you move each subprocess into its own process group, the `pgid` should be equal to the `pid`.

b. Foreground Terminal

Every terminal has an associated foreground process group ID. When you type Ctrl-C, your terminal sends a signal to every process inside the foreground process group. You can change which process group is in the foreground of a terminal with `tcsetpgrp(int fd, pid_t pgrp)`. The `fd` should be 0 for standard input `stdin`.

In your shell, you can use `kill -XXX pid`, where `XXX` is the human-friendly suffix of the desired signal, to send any signal to the process with process ID `pid`. Since you can use the signal function to change how signals are handled by the current process, your shell should basically ignore most of these signals, whereas your shell's subprocesses should respond with the default action. Be aware that forked processes will inherit the signal handlers of the original process. You may want to check out `man 2 signal` and `man 7 signal` for more information on this. You want to ensure that each program you start is in its own process group. When you start a process, its process group should be placed in the foreground. Stopping signals should only affect the foreground program(s), not the background shell.

DEFENSIVE PROGRAMMING

With regards to multiple commands on the same line, the following lines are all valid and have reasonable commands specified:

```
prompt> ls
prompt> /bin/ls
prompt> ls -l
prompt> ls -l; cat file1
prompt> ls -l | wc; cat file1
prompt> ls -l; cat file1; grep Dallas file2
```

For example, on the last line, the commands `ls -l`, `cat file`, and `grep Dallas file2` should all be running at the same time. As a result, you may see that their output

is intermixed. You must support pipes (i.e., `|`) to connect the output of one program to the input of another as in the second-to-last command above.

Check the return values of all functions utilizing system resources. Do not blindly assume all requests for memory will succeed and that all writes to a file will occur correctly. Your code should handle errors properly. Many failed function calls should not be fatal to a program. Typically, a system call will return `-1` in the case of an error (`malloc` returns `NULL` on error).

An OS cannot simply fail when it encounters an error. It must check all parameters before it trusts them. In general, there should be no circumstances in which your C program will core dump, hang indefinitely, or prematurely terminate. Therefore, your program must respond to all input in a reasonable manner. By “reasonable”, this means that you should print a meaningful and understandable error message and either continue processing or exit, depending upon the situation.

Many questions about functions and system behavior can be found in the manual pages.

You should consider the following situations as errors – in each case, your shell should print a message to `stderr` and exit gracefully:

- An incorrect number of command line arguments to your shell program; and
- The batch file does not exist or cannot be opened.

For the following situation, you should print a message to the user (`stderr`) and continue processing:

- A command does not exist or cannot be executed.

Optionally, to make coding your shell easier, you may print an error message and continue processing in the following situation:

- A very long command line (over 512 characters including the `'\n'`).

Your shell should also be able to handle the following scenarios, which are not errors (i.e., your shell should not print an error message):

- An empty command line;
- Extra white spaces within a command line; and
- Batch file ends without `exit` command or user types `'Ctrl-D'` as a command in interactive mode.

In no case should any input or any command-line format cause your shell program to crash or exit prematurely. You should think carefully about how you want to handle oddly formatted command lines (e.g., lines with no commands between a semi-colon). In these cases, you may choose to print a warning message and/or execute some subset of the commands. However, in all cases, your shell should continue to execute.

```
prompt> ; cat file1 ; grep Dallas file2
prompt> cat file1 ; ; grep Dallas file2
prompt> cat file1 ; ls -l ;
```

```
prompt> cat file1 ;;;; ls -l
prompt> ;; ls -l
prompt> ;
```

REQUIREMENTS

Your code must be written in C and be invoked exactly as follows:

```
newshell [batchFile]
```

The command-line arguments to your shell are to be interpreted as follows:

- `batchFile`: an optional argument (indicated by square brackets as above). If present, your shell will read each line of the `batchFile` for commands to be executed. If not present, your shell will run in interactive mode by printing a prompt to the user at `stdout` and reading the command `stdin`.

For example, if you run your program as:

```
newshell /home/mat0299/csce3600/batchfile
```

then it will read commands from `/home/mat0299/csce3600/batchfile` until it sees the `exit` command or EOF.

OPTIONAL SHELL FUNCTIONALITY

Teams who have completed all requirements for this program and are looking for an additional challenge may add the following optional functionality to gain bonus points added to your team's overall score:

- Add a shell history of previous commands run on the shell (+5 points).
- Allow the user to customize the prompt (+5 points).

However, all required functionality must be implemented prior to attempting this extra credit work as no points will be given for attempting this functionality if all requirements have not been completed. In other words, make sure your program is complete before attempting this extra credit.

GRADING

This assignment must be submitted via Canvas with the following elements:

- Your C program file(s). Your code should be well documented in terms of comments. For example, good comments in general consist of a header (with your name, course section, date, and brief description), comments for each variable, and commented blocks of code.
- A **README** file with some basic documentation about your code. This file should contain the following four components:
 - Your name(s).
 - Organization of the Project. Since there are multiple components in this project, you will describe how the work was organized and managed,

including which team members were responsible for what components – there are lots of ways to do this, so your team needs to come up with the best way that works based on your team's strengths. Note that this may be used in assessment of grades for this project.

- Design Overview: A few paragraphs describing the overall structure of your code and any important structures.
 - Complete Specification: Describe how you handled any ambiguities in the specification. For example, for this project, explain how your shell will handle lines that have no commands between semi-colons.
 - Known Bugs or Problems: A list of any features that you did not implement or that you know are not working correctly.
- A completed group assessment evaluation (given at a later date) for each team member. *Please be aware that a student receiving a poor evaluation with regards to their performance on the team will have his/her grading marks reduced by an appropriate amount, based on the evaluation.*
 - A `Makefile` for compiling your source code, including a clean directive.
 - Your program will be graded based largely on whether it works correctly on the CSE machines (e.g., `cse01`, `cse02`, ..., `cse06`), so you should make sure that your program compiles and runs on a CSE machine.

To ensure that your C code is compiled correctly, you will need to create a simple `Makefile`. This allows our scripts to just run `make` to compile your code with the right libraries and flags. When using `gcc` to compile your code, use the `-Wall` switch to ensure that all warnings are displayed. Do not be satisfied with code that merely compiles – it should compile with no warnings! You will lose points if your code produces warnings when compiled.

Here is a sample `Makefile` that may help as well:

```
#####  
#  
# shell.c is the name of your source code; you may change this.  
# However, you must keep the name of the executable as "shell".  
#  
# Type "make" or "make newshell" to compile your code  
#  
# Type "make clean" to remove the executable (and object files)  
#  
#####  
CC=gcc  
CFLAGS=-Wall -g  
newshell: shell.c  
    $(CC) -o newshell $(CFLAGS) shell.c  
clean:  
    $(RM) newshell
```

Your program will be tested using a suite of about 20 test cases on the CSE machines, some of which will exercise your program's ability to correctly execute commands and some of which will test your program's ability to catch error conditions. Be sure that you thoroughly exercise your program's capabilities on a wide range of test suites.

SUBMISSION

- Each team will electronically submit all components to the **Major 1** dropbox in Canvas by the due date. Group submission will be enabled on Canvas so that only one team member needs to submit the assignment.